

Retrieval-Augmented Generation (RAG) is an AI technique that combines document retrieval with powerful language models. It improves accuracy by grounding responses in real data. LangChain provides tools for embeddings, vector stores, chains, and prompts that make RAG easy to build. LangGraph helps orchestrate decision-making, such as switching between weather API calls or document-based answers. FAISS is a fast vector database used to index and search document chunks. Together, these tools allow developers to build scalable, accurate, and modular AI applications.